

# Neural Networks

## 1 Introduction

The objective of this practical is to get a grasp of Neural Networks (NN) with Python. The reader should refer to the lecture slides of the course in “Artificial Intelligence”.

Given a neural network characterized by 2 inputs, 2 outputs and a hidden layer, you will proceed as follows:

1. With an initial guess of the weights, implement a python function for doing a forward pass of the network (Sect. 2).
2. In the same function, implement the backpropagation of the network, and test everything with a dataset of two labeled samples. (Sect. 2).
3. With the same dataset, run the network (forward and back) to train it over *many* iterations (Sect. 3.1).
4. Finally, train and then test the network with the bigger dataset provided in with the provided datasets `NNTraining.data` and `NNTest.data` (Sect. 3.2).

For this practical, you will need some python packages. Add the following lines to your python script to ensure that all packages are imported correctly.

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

## 2 Programming the network

The goal of this section is to design the network shown in Fig. 1.

You dispose of two labeled samples. Each sample contains input  $\mathbf{x}$  with corresponding target output  $\mathbf{t}$ :

$$\mathbf{x}_1 = \begin{bmatrix} 2 \\ 1 \end{bmatrix} \quad \mathbf{t}_1 = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \mathbf{x}_2 = \begin{bmatrix} -1 \\ 3 \end{bmatrix} \quad \mathbf{t}_2 = \begin{bmatrix} 0 \\ 1 \end{bmatrix} \quad (1)$$

Proceed as follows.

1. Design a function for the forward pass, which takes as arguments vectors  $\mathbf{x}$ ,  $\mathbf{t}$  and  $\mathbf{w}$  and returns the loss  $L$  (squared error between  $\mathbf{y}$  and  $\mathbf{t}$ ). Initialize the network with the following weights:

$$\mathbf{w} = [2 \quad -3 \quad -3 \quad 4 \quad 1 \quad -1 \quad 0.25 \quad 2]^\top. \quad (2)$$

You should obtain:  $L = 0.245$  and  $L = 0.330$ , respectively. For this exercise, it is wise to define a function returning the sigmoid output  $\sigma(z)$  (see lecture slides).

2. In the same function, implement the backpropagation of the network. Along with the loss  $L$ , the function should now return its gradient  $\nabla_{\mathbf{w}} L$ . You should obtain:

$$\begin{aligned} \nabla_{\mathbf{w}} L &= [-0.12 \quad 0.07 \quad -0.06 \quad 0.03 \quad -0.10 \quad 0.13 \quad -0.02 \quad 0.02]^\top \\ \nabla_{\mathbf{w}} L &= [\sim 0 \quad \sim 0 \quad \sim 0 \quad \sim 0 \quad \sim 0 \quad \sim 0 \quad 0.28 \quad -0.025]^\top. \end{aligned} \quad (3)$$




Figure 1: Neural network studied in this practical.

## 3 Training and testing the network

### 3.1 Training and testing on three samples

With the same dataset, run the network (forward and back) to train it until the loss function has diminished *enough*. Since the dataset consists in only two samples, we use a single batch coinciding with the whole dataset.

You will need to manually tune two hyperparameters:

1. the gradient descent step size  $\rho$ ,
2. the number of times (iterations)  $N$  you compute the mean of the gradient over the batch<sup>1</sup>.

After each iteration, store the loss  $L$  so you can plot its trend over the number of iterations. Tune  $\rho$  and  $N$  until you find a satisfactory trend (as in Fig. 2).

Test with a new labeled sample:

$$\mathbf{x}_3 = \begin{bmatrix} 1 \\ 4 \end{bmatrix} \quad \mathbf{t}_3 = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad (4)$$

and comment.

---

<sup>1</sup>Here, we use this criterion instead of the tolerance on weight variation seen in class.




Figure 2: Evolution of the loss when training with two samples.

### 3.2 Training and testing on a large dataset

Now, train the network on the provided dataset `NNTraining.data`, starting from the same weight as before (equation (2)). To load the dataset you can use the following snippet:

```
data = pd.read_csv('NNTraining.data')
x1 = data.iloc[:,0]
x2 = data.iloc[:,1]
t1 = data.iloc[:,2]
t2 = data.iloc[:,3]
```

Since this dataset is much larger (200 samples) it can be split in equally sized batches for more efficient learning. After all samples in a batch have been passed forward, compute the average gradient over all samples, and update  $\mathbf{w}$ . After each epoch (i.e., each time all batches have been passed forward), store loss  $L$  so you can plot its trend at the end of training.

This time you will need to manually tune three hyperparameters:

1. the gradient descent step size  $\rho$ ,
2. the batch size  $m$ ,
3. the number of epochs  $N$ .

Start with a single batch of size  $m = 200$ , then increase the number of batches and check how the loss function evolves. After a bit of hyperparameter tuning, I obtained the curve in Fig. 3. Compare the final loss you obtain here with the one obtained in Sec. 3.1.




Figure 3: Evolution of the loss when training with `NNTraining.data`.

How does the trained network generalize to new data? To assess this, test the network (only forward) on the test dataset `NNTest.data` and plot the 100 losses. You can plot them before and after training to see the difference, as I did in Fig. 4.




Figure 4: Loss obtained when passing each of the 100 samples of `NNTest.data` through the network.